



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

Experiment No. 6
Apply Boosting Algorithm on Adult Census Income Dataset and analyze the performance of the model
Date of Performance: 04-09-2023
Date of Submission: 13-09-2023



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

Aim: Apply Boosting algorithm on Adult Census Income Dataset and analyze the performance of the model.

Objective: Apply Boosting algorithm on the given dataset and maximize the accuracy, Precision, Recall, F1 score.

Theory:

Suppose that as a patient, you have certain symptoms. Instead of consulting one doctor, you choose to consult several. Suppose you assign weights to the value or worth of each doctor's diagnosis, based on the accuracies of previous diagnosis they have made. The final diagnosis is then a combination of the weighted diagnosis. This is the essence behind boosting.

Algorithm: Adaboost- A boosting algorithm—creates an ensemble of classifiers. Each one gives a weighted vote.

Input:

- D , a set of d class labelled training tuples
- k , the number of rounds (one classifier is generated per round)
- a classification learning scheme

Output: A composite model

Method

1. Initialize the weight of each tuple in D is $1/d$
2. For $i=1$ to k do // for each round
3. Sample D with replacement according to the tuple weights to obtain D_i
4. Use training set D_i to derive a model M_i
5. Compute $\text{error}(M_i)$, the error rate of M_i
6. $\text{Error}(M_i) = \sum w_j * \text{err}(X_j)$
7. If $\text{Error}(M_i) > 0.5$ then
8. Go back to step 3 and try again
9. endif
10. for each tuple in D_i that was correctly classified do
11. Multiply the weight of the tuple by $\text{error}(M_i)/(1-\text{error}(M_i))$
12. Normalize the weight of each tuple
13. end for



To use the ensemble to classify tuple X

1. Initialize the weight of each class to 0
2. for $i=1$ to k do // for each classifier
3. $w_i = \log((1 - \text{error}(M_i)) / \text{error}(M_i))$ // weight of the classifiers vote
4. $C = M_i(X)$ // get class prediction for X from M_i
5. Add w_i to weight for class C
6. end for
7. Return the class with the largest weight.

Dataset:

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:

>50K, <=50K.

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.

capital-gain: continuous.

capital-loss: continuous.

hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad & Tobago, Peru, Hong, Holand-Netherlands.

12-aditi-sawant-ml-exp6

September 12, 2023

```
[ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import io
from sklearn.metrics import accuracy_score, precision_score, f1_score, \
    confusion_matrix, classification_report
from sklearn.model_selection import cross_val_score
from sklearn.metrics import mean_squared_error
```

```
[ ]: import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

```
[ ]: file = ('/content/adult.csv')
df = pd.read_csv(file)
```

```
[ ]: print(df.head())
```

	age	workclass	fnlwgt	education	education.num	marital.status	\
0	90	?	77053	HS-grad	9	Widowed	
1	82	Private	132870	HS-grad	9	Widowed	
2	66	?	186061	Some-college	10	Widowed	
3	54	Private	140359	7th-8th	4	Divorced	
4	41	Private	264663	Some-college	10	Separated	

	occupation	relationship	race	sex	capital.gain	\
0	?	Not-in-family	White	Female	0	
1	Exec-managerial	Not-in-family	White	Female	0	
2	?	Unmarried	Black	Female	0	
3	Machine-op-inspct	Unmarried	White	Female	0	
4	Prof-specialty	Own-child	White	Female	0	

	capital.loss	hours.per.week	native.country	income
0	4356	40	United-States	<=50K
1	4356	18	United-States	<=50K
2	4356	40	United-States	<=50K

3	3900	40	United-States	<=50K
4	3900	40	United-States	<=50K

```
[ ]: print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   32561 non-null  int64
1   workclass             32561 non-null  object
2   fnlwgt               32561 non-null  int64
3   education             32561 non-null  object
4   education.num        32561 non-null  int64
5   marital.status       32561 non-null  object
6   occupation            32561 non-null  object
7   relationship         32561 non-null  object
8   race                 32561 non-null  object
9   sex                  32561 non-null  object
10  capital.gain          32561 non-null  int64
11  capital.loss          32561 non-null  int64
12  hours.per.week       32561 non-null  int64
13  native.country       32561 non-null  object
14  income               32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
None
```

```
[ ]: for i in df.columns:
      t = df[i].value_counts()
      index = list(t.index)
      print ("Count of ? in", i)
      for i in index:
          temp = 0
          if i == '?':
              print (t['?'])
              temp = 1
              break
      if temp == 0:
          print ("0")
```

Count of ? in age

0

Count of ? in workclass

1836

Count of ? in fnlwgt

0

```

Count of ? in education
0
Count of ? in education.num
0
Count of ? in marital.status
0
Count of ? in occupation
1843
Count of ? in relationship
0
Count of ? in race
0
Count of ? in sex
0
Count of ? in capital.gain
0
Count of ? in capital.loss
0
Count of ? in hours.per.week
0
Count of ? in native.country
583
Count of ? in income
0

```

```
[ ]: df=df.loc[(df['workclass'] != '?') & (df['native.country'] != '?')]
print(df.head())
```

	age	workclass	fnlwgt	education	education.num	marital.status	\
1	82	Private	132870	HS-grad	9	Widowed	
3	54	Private	140359	7th-8th	4	Divorced	
4	41	Private	264663	Some-college	10	Separated	
5	34	Private	216864	HS-grad	9	Divorced	
6	38	Private	150601	10th	6	Separated	

	occupation	relationship	race	sex	capital.gain	\
1	Exec-managerial	Not-in-family	White	Female	0	
3	Machine-op-inspct	Unmarried	White	Female	0	
4	Prof-specialty	Own-child	White	Female	0	
5	Other-service	Unmarried	White	Female	0	
6	Adm-clerical	Unmarried	White	Male	0	

	capital.loss	hours.per.week	native.country	income
1	4356	18	United-States	<=50K
3	3900	40	United-States	<=50K
4	3900	40	United-States	<=50K
5	3770	45	United-States	<=50K
6	3770	40	United-States	<=50K

```
[ ]: df["income"] = [1 if i=='>50K' else 0 for i in df["income"]]
print(df.head())
```

	age	workclass	fnlwgt	education	education.num	marital.status	\
1	82	Private	132870	HS-grad	9	Widowed	
3	54	Private	140359	7th-8th	4	Divorced	
4	41	Private	264663	Some-college	10	Separated	
5	34	Private	216864	HS-grad	9	Divorced	
6	38	Private	150601	10th	6	Separated	

	occupation	relationship	race	sex	capital.gain	\
1	Exec-managerial	Not-in-family	White	Female	0	
3	Machine-op-inspct	Unmarried	White	Female	0	
4	Prof-specialty	Own-child	White	Female	0	
5	Other-service	Unmarried	White	Female	0	
6	Adm-clerical	Unmarried	White	Male	0	

	capital.loss	hours.per.week	native.country	income
1	4356	18	United-States	0
3	3900	40	United-States	0
4	3900	40	United-States	0
5	3770	45	United-States	0
6	3770	40	United-States	0

```
[ ]: df_more=df.loc[df['income'] == 1]
print(df_more.head())
```

	age	workclass	fnlwgt	education	education.num	marital.status	\
7	74	State-gov	88638	Doctorate	16	Never-married	
10	45	Private	172274	Doctorate	16	Divorced	
11	38	Self-emp-not-inc	164526	Prof-school	15	Never-married	
12	52	Private	129177	Bachelors	13	Widowed	
13	32	Private	136204	Masters	14	Separated	


	occupation	relationship	race	sex	capital.gain	\
7	Prof-specialty	Other-relative	White	Female	0	
10	Prof-specialty	Unmarried	Black	Female	0	
11	Prof-specialty	Not-in-family	White	Male	0	
12	Other-service	Not-in-family	White	Female	0	
13	Exec-managerial	Not-in-family	White	Male	0	

	capital.loss	hours.per.week	native.country	income
7	3683	20	United-States	1
10	3004	35	United-States	1
11	2824	45	United-States	1
12	2824	20	United-States	1
13	2824	55	United-States	1



```
[ ]: workclass_types = df_more['workclass'].value_counts()
labels = list(workclass_types.index)
aggregate = list(workclass_types)
print(workclass_types)
print(aggregate)
print(labels)
```

```
Private          4876
Self-emp-not-inc  714
Local-gov        609
Self-emp-inc     600
Federal-gov      365
State-gov        344
Name: workclass, dtype: int64
[4876, 714, 609, 600, 365, 344]
['Private', 'Self-emp-not-inc', 'Local-gov', 'Self-emp-inc', 'Federal-gov',
'State-gov']
```

```
[ ]: plt.figure(figsize=(7,7))
plt.pie(aggregate, labels=labels, autopct='%1.1f%%', shadow = True)
plt.axis('equal')
plt.show()
```



```
[ ]: #Count plot on single categorical variable
sns.countplot(x='income', data = df)
plt.show()
df['income'].value_counts()
```




```
[ ]: 0    22661  
     1     7508  
     Name: income, dtype: int64
```

```
[ ]: #Plot figsize  
plt.figure(figsize=(10,7))  
sns.heatmap(df.corr(), cmap='coolwarm', annot=True)  
print(plt.show())
```

<ipython-input-13-6201d8194dba>:3: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
sns.heatmap(df.corr(), cmap='coolwarm', annot=True)
```



None

```
[ ]: plt.figure(figsize=(10,7))
sns.distplot(df['age'], color="red", bins=100)
plt.ylabel("Distribution", fontsize = 10)
plt.xlabel("Age", fontsize = 10)
plt.show()
```


<ipython-input-14-1b72b8b67fa9>:2: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['age'], color="red", bins=100)
```



```
[ ]: #To find distribution of categorical columns w.r.t income
fig, axes = plt.subplots(figsize=(20, 10))

plt.subplot(231)
sns.countplot(x='workclass',
              hue='income',
              data = df,
              palette="BuPu")
plt.xticks(rotation=90)

plt.subplot(232)
sns.countplot(x='marital.status',
              hue='income',
              data = df,
              palette="deep")
plt.xticks(rotation=90)

plt.subplot(233)
sns.countplot(x='education',
              hue='income',
              data = df,
              palette = "autumn")
```

```

plt.xticks(rotation=90)

plt.subplot(234)
sns.countplot(x='relationship',
              hue='income',
              data=df,
              palette="inferno")
plt.xticks(rotation=90)

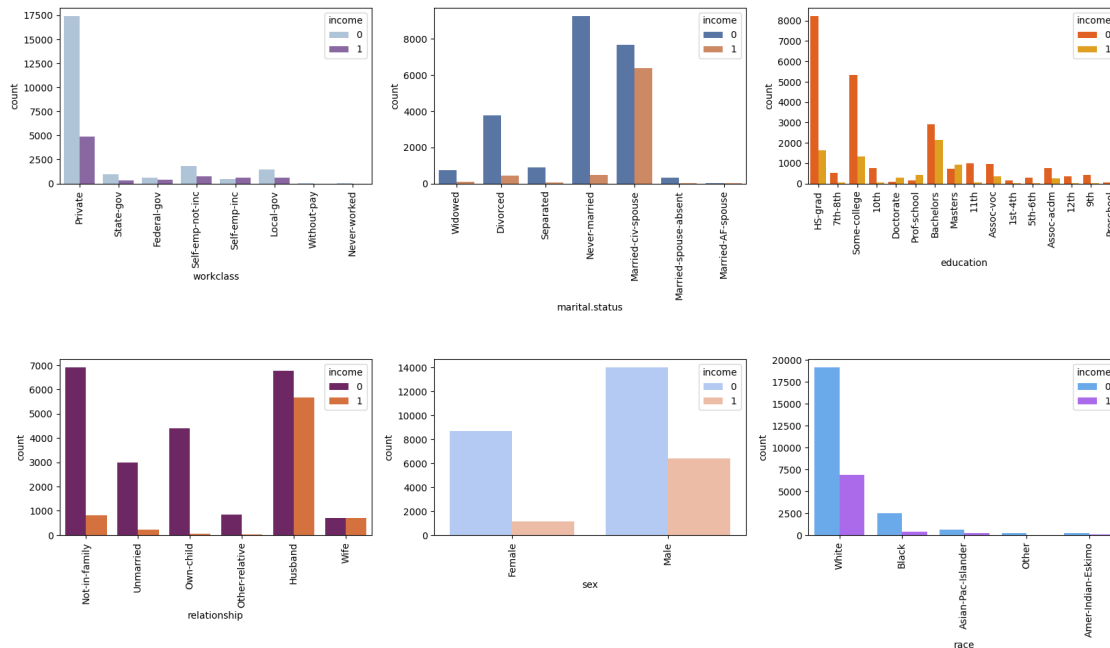
plt.subplot(235)
sns.countplot(x='sex',
              hue='income',
              data=df,
              palette="coolwarm")
plt.xticks(rotation=90)

plt.subplot(236)
sns.countplot(x='race',
              hue='income',
              data=df,
              palette="cool")
plt.xticks(rotation=90)
plt.subplots_adjust(hspace=1)
plt.show()

```

<ipython-input-15-f6a96c604872>:4: MatplotlibDeprecationWarning: Auto-removal of overlapping axes is deprecated since 3.6 and will be removed two minor releases later; explicitly call ax.remove() as needed.

```
plt.subplot(231)
```



```
[ ]: df1 = df.copy()
```

```
[ ]: categorical_features = list(df1.select_dtypes(include=['object']).columns)
print(categorical_features)
df1
```

```
['workclass', 'education', 'marital.status', 'occupation', 'relationship',
'race', 'sex', 'native.country']
```

```
[ ]:
    age workclass  fnlwgt    education  education.num    marital.status \
1      82   Private 132870      HS-grad             9      Widowed
3      54   Private 140359      7th-8th             4      Divorced
4      41   Private 264663  Some-college            10      Separated
5      34   Private 216864      HS-grad             9      Divorced
6      38   Private 150601       10th              6      Separated
...
32556  22   Private 310152  Some-college            10      Never-married
32557  27   Private 257302   Assoc-acdm            12  Married-civ-spouse
32558  40   Private 154374      HS-grad             9  Married-civ-spouse
32559  58   Private 151910      HS-grad             9      Widowed
32560  22   Private 201490      HS-grad             9      Never-married

    occupation  relationship  race  sex  capital.gain \
1   Exec-managerial  Not-in-family  White  Female      0
3  Machine-op-inspct   Unmarried  White  Female      0
4   Prof-specialty    Own-child  White  Female      0
```

5	Other-service	Unmarried	White	Female	0
6	Adm-clerical	Unmarried	White	Male	0
...
32556	Protective-serv	Not-in-family	White	Male	0
32557	Tech-support	Wife	White	Female	0
32558	Machine-op-inspct	Husband	White	Male	0
32559	Adm-clerical	Unmarried	White	Female	0
32560	Adm-clerical	Own-child	White	Male	0

	capital.loss	hours.per.week	native.country	income
1	4356	18	United-States	0
3	3900	40	United-States	0
4	3900	40	United-States	0
5	3770	45	United-States	0
6	3770	40	United-States	0
...
32556	0	40	United-States	0
32557	0	38	United-States	0
32558	0	40	United-States	1
32559	0	40	United-States	0
32560	0	20	United-States	0

[30169 rows x 15 columns]

```
[ ]: from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
for feat in categorical_features:
    df1[feat] = le.fit_transform(df1[feat].astype(str))
df1
```

[]:	age	workclass	fnlwgt	education	education.num	marital.status	\
1	82	3	132870	11	9	6	
3	54	3	140359	5	4	0	
4	41	3	264663	15	10	5	
5	34	3	216864	11	9	0	
6	38	3	150601	0	6	5	
...	
32556	22	3	310152	15	10	4	
32557	27	3	257302	7	12	2	
32558	40	3	154374	11	9	2	
32559	58	3	151910	11	9	6	
32560	22	3	201490	11	9	4	

	occupation	relationship	race	sex	capital.gain	capital.loss	\
1	4	1	4	0	0	4356	
3	7	4	4	0	0	3900	
4	10	3	4	0	0	3900	

5	8	4	4	0	0	3770
6	1	4	4	1	0	3770
...
32556	11	1	4	1	0	0
32557	13	5	4	0	0	0
32558	7	0	4	1	0	0
32559	1	4	4	0	0	0
32560	1	3	4	1	0	0

	hours.per.week	native.country	income
1	18	38	0
3	40	38	0
4	40	38	0
5	45	38	0
6	40	38	0
...
32556	40	38	0
32557	38	38	0
32558	40	38	1
32559	40	38	0
32560	20	38	0

[30169 rows x 15 columns]

```
[ ]: X = df1.drop(columns = ['income'])
y = df1['income'].values

# Splitting the data set into train and test set
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.
↳3, random_state = 0)

print ("Train set size: ", X_train.shape)
print ("Test set size: ", X_test.shape)
```

Train set size: (21118, 14)
Test set size: (9051, 14)

```
[ ]: from sklearn.ensemble import AdaBoostClassifier

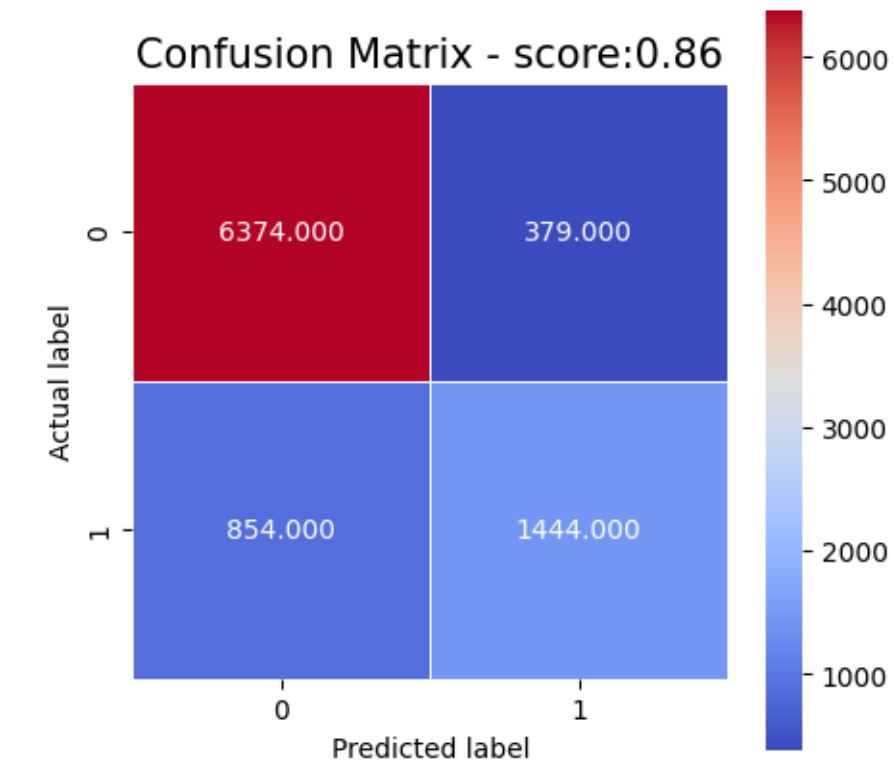
# Train Adaboost Classifier
abc = AdaBoostClassifier(n_estimators = 300, learning_rate=1)
abc_model = abc.fit(X_train, y_train)

#Prediction
y_pred_abc = abc_model.predict(X_test)
```

```
print("Accuracy: ", accuracy_score(y_test, y_pred_abc))
print("F1 score :", f1_score(y_test, y_pred_abc, average='binary'))
print("Precision : ", precision_score(y_test, y_pred_abc))
```

Accuracy: 0.8637719588995691
 F1 score : 0.7008007765105557
 Precision : 0.7921009325287987

```
[ ]: cm = confusion_matrix(y_test, y_pred_abc)
plt.figure(figsize=(5,5))
sns.heatmap(cm, annot=True, fmt=".3f", linewidths=.5, square = True, cmap = 'coolwarm');
plt.ylabel('Actual label');
plt.xlabel('Predicted label');
plt.title('Confusion Matrix - score:' + str(round(accuracy_score(y_test, y_pred_abc), 2)), size = 15);
plt.show()
print(classification_report(y_test, y_pred_abc))
```



	precision	recall	f1-score	support
0	0.88	0.94	0.91	6753

1	0.79	0.63	0.70	2298
accuracy			0.86	9051
macro avg	0.84	0.79	0.81	9051
weighted avg	0.86	0.86	0.86	9051

```
[ ]: from sklearn.ensemble import GradientBoostingClassifier

#Training the model with gradient boosting
gbc = GradientBoostingClassifier(
    learning_rate = 0.1,
    n_estimators = 500,
    max_depth = 5,
    subsample = 0.9,
    min_samples_split = 100,
    max_features='sqrt',
    random_state=10)
gbc.fit(X_train,y_train)

# Predictions
y_pred_gbc = gbc.predict(X_test)


print("Accuracy : ",accuracy_score(y_test, y_pred_gbc))
print("F1 score : ", f1_score(y_test, y_pred_gbc, average = 'binary'))
print("Precision : ", precision_score(y_test, y_pred_gbc))
```

```
Accuracy : 0.8689647552756602
F1 score : 0.7218574108818011
Precision : 0.7828077314343845
```

```
[ ]: rms = np.sqrt(mean_squared_error(y_test, y_pred_gbc))
print("RMSE for gradient boost: ", rms)
```

```
RMSE for gradient boost: 0.3619879068758235
```

```
[ ]: cm = confusion_matrix(y_test, y_pred_gbc)
plt.figure(figsize=(5,5))
sns.heatmap(cm, annot = True, fmt=".3f", linewidths = 0.5, square = True, cmap=
    ⇨ "coolwarm");
plt.ylabel('Actual label');
plt.xlabel('Predicted label');
plt.title('Confusion Matrix - score:' + str(round(accuracy_score(y_test,
    ⇨ y_pred_gbc),2)), size = 15);
plt.show()
print(classification_report(y_test, y_pred_gbc))
```



	precision	recall	f1-score	support
0	0.89	0.94	0.91	6753
1	0.78	0.67	0.72	2298
accuracy			0.87	9051
macro avg	0.84	0.80	0.82	9051
weighted avg	0.86	0.87	0.87	9051

```
[ ]: import xgboost as xgb
from xgboost import XGBClassifier

#Training the model with gradient boosting
xgboost = XGBClassifier(learning_rate=0.01,
                        colsample_bytree = 0.4,
                        n_estimators=1000,
                        max_depth=20,
                        gamma=1)

xgboost_model = xgboost.fit(X_train, y_train)
```

```
# Predictions
y_pred_xgboost = xgboost_model.predict(X_test)


print("Accuracy : ", accuracy_score(y_test, y_pred_xgboost))
print("F1 score : ", f1_score(y_test, y_pred_xgboost, average = 'binary'))
print("Precision : ", precision_score(y_test, y_pred_xgboost))
```

```
Accuracy : 0.8655397193680257
F1 score : 0.7090604829070045
Precision : 0.786737400530504
```

```
[ ]: rms = np.sqrt(mean_squared_error(y_test, y_pred_xgboost))
print("RMSE for xgboost: ", rms)
```


```
RMSE for xgboost: 0.3666882608319693
```

```
[ ]: cm = confusion_matrix(y_test, y_pred_xgboost)
plt.figure(figsize=(5,5))
sns.heatmap(cm, annot=True, fmt=".3f", linewidths=.5, square = True, cmap = "coolwarm");
plt.ylabel('Actual label');
plt.xlabel('Predicted label');
plt.title('Confusion Matrix - score: '+str(round(accuracy_score(y_test, y_pred_xgboost),2)), size = 15);
plt.show()
print(classification_report(y_test, y_pred_xgboost))
```



	precision	recall	f1-score	support
0	0.89	0.94	0.91	6753
1	0.79	0.65	0.71	2298
accuracy			0.87	9051
macro avg	0.84	0.79	0.81	9051
weighted avg	0.86	0.87	0.86	9051

```
[ ]: from sklearn.metrics import roc_curve
fpr, tpr, thresholds = roc_curve(y_test, xgboost.predict_proba(X_test)[: ,1])
plt.figure(figsize = (10,5))
plt.plot([0,1],[0,1], 'k--')
plt.plot(fpr, tpr)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC CURVE for Xgboost')
plt.show()
```





Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

Conclusion:

1. The GradientBoostingClassifier has the highest accuracy (0.8690) and F1 score (0.7219) among the three classifiers, indicating that it performs the best on this dataset.
2. The AdaBoostClassifier and XGBClassifier also have good performance but slightly lower than the GradientBoostingClassifier in terms of accuracy and F1 score.
3. All three classifiers have relatively high precision, indicating that when they predict a positive class (1), they are often correct.
4. The recall values vary among the classifiers, with the GradientBoostingClassifier having the highest recall for the positive class (1), indicating that it correctly identifies more positive cases.
5. The F1 score, which balances precision and recall, shows how well the classifiers perform in classifying both classes.

Comparison between boosting algorithms and random forest classifier:

1. The boosting algorithms (AdaBoost, Gradient Boosting, and XGBoost) generally outperform the Random Forest Classifier in terms of accuracy, precision, and F1 score.
2. The Random Forest Classifier performs reasonably well, with an accuracy of around 85% and a balanced F1 score. However, it falls slightly behind the boosting algorithms.
3. The boosting algorithms tend to have higher precision and recall for the positive class (income > 50K), indicating better ability to correctly classify high-income individuals.

All these models perform reasonably well on the Adult Census Income Dataset; the boosting algorithms, especially Gradient Boosting, appear to provide slightly better results in terms of accuracy and F1 score compared to the Random Forest Classifier.